Chapter 8: Relational Algebra

Outline:

- Introduction
- Unary Relational Operations.
    - Select Operator (σ)
    - Project Operator (π)
    - Rename Operator (ρ)
    - Assignment Operator (←)

- Binary Relational Operations.
    - Set Operators
        - Union Operator (∪)
        - Intersection Operator (∩)
        - Set Difference or Minus Operator (-)
    - Cartesian Product Operator (×)
    - Join Operator
        - Theta Join ($\bowtie_{\square}$)
        - Natural Join (⋈) or (*)

- Examples of Queries in Relational Algebra.

8.1 Introduction

Relational Query Languages

- Languages for describing queries on a relational database
- Structured Query Language (SQL)
  - Declarative (Nonprocedural)
- Relational Algebra
  - Intermediate language used within DBMS
  - Procedural
- Procedural: Relational expression specifies query by describing an algorithm (the sequence in which operators are applied) for determining the result of an expression

- Relational algebra is the basic set of operations for the relational model
- These operations enable a user to specify basic retrieval requests (or queries)
- The result of an operation is a new relation, which may have been formed from one or more input relations
  - This property makes the algebra "closed" (all objects in relational algebra are relations)
- The algebra operations thus produce new relations
  - These can be further manipulated using operations of the same algebra
- A sequence of relational algebra operations forms a relational algebra expression
  - The result of a relational algebra expression is also a relation that represents the result of a database query (or retrieval request)
- The fundamental operations in the relational algebra are select, project, union, set difference, cartesian product, and rename.

The Role of Relational Algebra in a DBMS

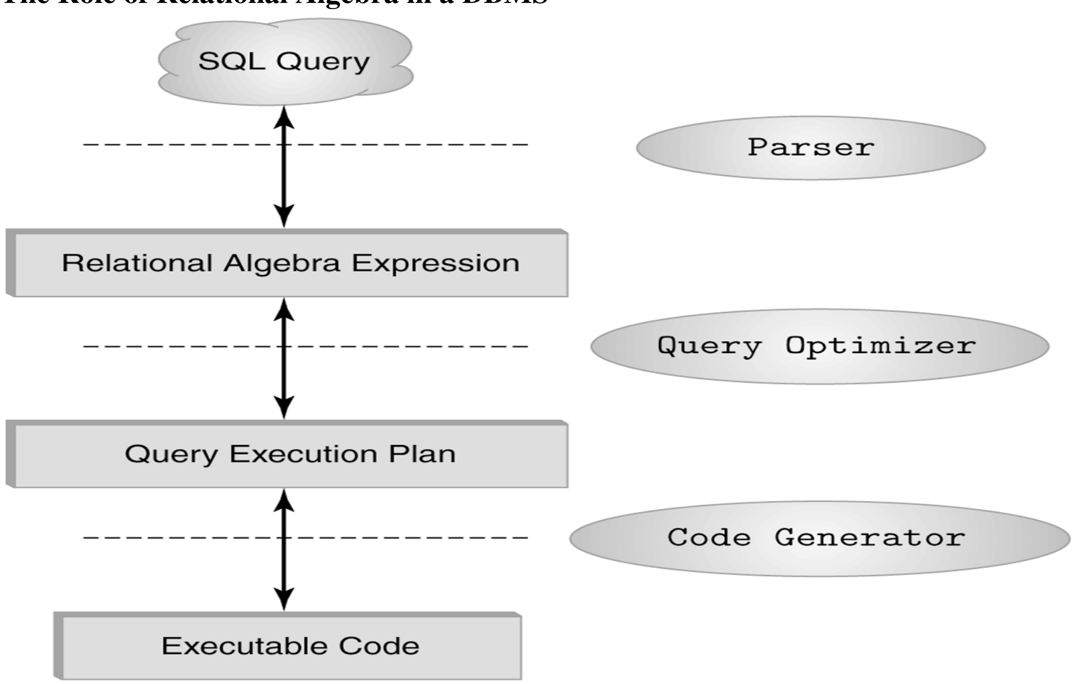

8.2 Unary Relational Operations

Select Operator (σ)

- Select a subset of rows from a relation that satisfying a condition.
  Syntax:

  $\sigma_{condition}(R)$

  - The symbol σ (sigma) is used to denote the select operator.
  - The selection condition is a Boolean expression specified on the attributes of relation R.
- $\sigma_{condition}(R)$ is equivalent to "Select * from R where <condition>;"

- Example: Consider the following relation r.

r

A	B	C	D
X	X	1	8
X	Y	5	7
Y	Y	3	3
Y	Y	12	10

١. $\sigma_{A=B}(r)$ = Select * from r where A=B;

A	B	C	D
X	X	1	8
Y	Y	3	3
Y	Y	12	10

٢. $\sigma_{D>5}(r)$ = Select * from r where D>5;

A	B	C	D
X	X	1	8
X	Y	5	7
Y	Y	12	10

- We can combine several conditions into a larger condition by using the connectives ^ (and), □ (or), and ¬ (not).
- Example:

  $\sigma_{A=B \wedge D>5}(r)$ = Select * from r where A=B and D>5;

A	B	C	D
X	X	1	8
Y	Y	12	10

- Example: Retrieve the Id, Name, Address of students who live in Amman.

Student

Id	Name	Address
1108	Ali	Amman
1453	Ahmad	Salt
1002	Omar	Amman
2603	Anas	Zarqa

$\sigma_{address='Amman'}$ (Student) = Select * from Student where address='Amman';

Id	Name	Address
1108	Ali	Amman
1002	Omar	Amman

- Example: Retrieve the Id, Name, Address of student who's name is Ahmad or students who live in Amman.

  $\sigma_{name='Ahmad' \lor address='Amman'}$ (Student) =
  Select * from Student where name = 'Ahmad' or address = 'Amman';

Id	Name	Address
1108	Ali	Amman
1453	Ahmad	Salt
1002	Omar	Amman

- Select Operation Properties:
  $\sigma_{condition1}(\sigma_{condition2}(R)) = \sigma_{condition2}(\sigma_{condition1}(R)) = \sigma_{condition1 \wedge condition2}(R)$

Project Operator (π)

- Selecting a subset of the attributes of a relation by specifying the name of the required attributes.
- The Project creates a vertical partitioning.
  Syntax:

  $\pi_{\text{<Attribute list>}}(R)$

  - The symbol π (pi) is used to denote the project operator.
  - <Attribute list> is the desired list of attributes from the attributes of relation R.
  - $\pi_{\text{<Attribute list>}}(R)$ is equivalent to "Select Distinct *Attribute_List* from R;"
- The project operation removes any duplicate rows.
- Example: Consider the following relation r.

r

A	B	C	D
X	X	1	8
X	Y	5	7
Y	Y	3	3
Y	Y	12	10

$\pi_{A,B}(r)$ = Select Distinct A, B From r;

A	B
X	X
X	Y
Y	Y

- Project operation properties:
  - $\pi_{list1}(\pi_{list2}(r)) = \pi_{list1}(r)$, where list2 contains the attributes of list1
  - The number of rows in the result of projection $\pi_{list}(r)$ is always less or equal to the number of rows in r.
  - If the list of attributes includes a key of r, the number of rows is equal to the number of rows in r.

- Composition of Relational Operations (Expression)
  - Relational algebra operations can be composed together into a relational algebra expression.
  - Example:

    $\pi_B(\sigma_{C>=3}(r))$ = Select Distinct B From r Where C >= 3;

B
Y

Assignment Operator (⟵)

- We may want to apply several relational algebra operations one after other. Either we can write the operations as a single relational algebra expression by nesting the operations, or we can apply one operation at a time and create intermediate result relations. In the latter case, we must give names to the relations that hold the intermediate results.
- Example:

  $r1 \leftarrow \sigma_{C>=3}(r)$

  $\pi_B(r1)$

B
Y

Rename Operator (ρ)

- Allows us to refer to a relation by more than one name.
  Syntax:

  ١. ρ □ (□)

  Returns the expression R under the name X

  ٢. ρ □ (□□, □□, .. ,□□) (□)

If a relational-algebra expression R has arity n, then returns the result of expression R under the name X, and with the attributes renamed to □□ , □□ , …., □□ .

٣. ρ (□□, □□, .. ,□□) (□)

Rename the attributes names without changing the relation name

8.3 Binary Relational Operations

Set Operators

- Relation is a set of tuples, so set operations should apply: $\cap$, $\cup$, $-$ (set difference)
- Result of combining two relations with a set operator is a relation (all its elements must be tuples having same structure).

١. **Union Operator** ($\cup$)

- The result of this operation, denoted by R $\cup$ S, is a relation that includes all rows that are either in R or in S or in both R and S. Duplicate rows are eliminated.
- Example:

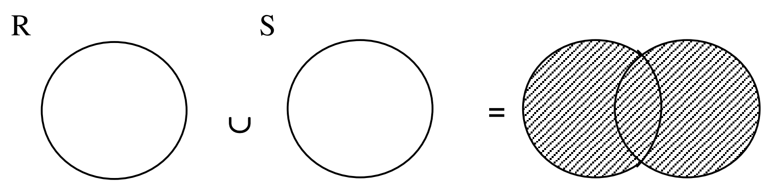

- For R $\cup$ S to be valid. (The two operands must be “type compatible”)
  ١. R, S must have the same arity (same number of attributes).
  ٢. The attribute domains must be compatible (example: 2nd column of R deals with the same type of values as does the 2nd column of S)

- R $\cup$ S is equivalent to “Select * From R Union Select * From S;”
- Example:

Tables:

Person (*SSN, Name, Address, Hobby*)
Professor (*Id, Name, Office, Phone*)

are **<u>not</u>** union compatible.
But

π_{Name} (Person) and π_{Name} (Professor)

<u>are</u> union compatible so

π_{Name} (Person) $\cup$ π_{Name} (Professor)

makes sense.

- Example: **r**

A	B
X	1
X	2
Y	1

s

A	B
X	2
Y	3

$r \cup s$ = Select * From r Union Select * from s;

A	B
X	1
X	2
Y	1
Y	3

- Example: retrieve the SSNs of all employees who either work in department 5 or directly supervise an employee who works in department 5.

Employee

<u>SSN</u>	EName	Sal	SuperSSN	DNo

$Dep5\_Emps \leftarrow \sigma_{DNo=5}\,(Employee)$
$Result1 \leftarrow \pi_{SSN}\,(Dep5\_Emps)$
$Result2 \leftarrow \pi_{SuperSSN}\,(Dep5\_Emps)$
$Result \leftarrow Result1 \cup Result2$

٢. **Intersection Operator** (∩)

- The result of this operation, denoted by R ∩ S, is a relation that includes all rows that are in both R and S. the two operands must be "type compatible".
- Example:

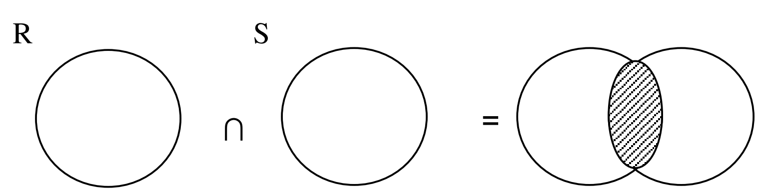

- R ∩ S is equivalent to "Select * From R Intersect Select * From S;"

- Example:

r

A	B
X	1
X	2

Y	1

s

A	B
X	2
Y	3

r ∩ s = Select * from r Intersect Select * from s;

A	B
X	2

٣. Intersection Operator (-)

- The result of this operation, denoted by R - S, is a relation that includes all rows that are in R but not in S. the two operands must be "type compatible".
- Example:

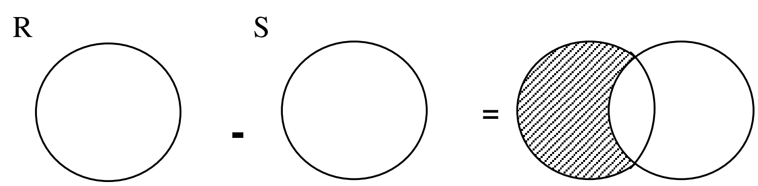

- R - S is equivalent to "Select * From R Minus Select * From S;"
- Example:

r

A	B
X	1
X	2
Y	1

s

A	B
X	2
Y	3

r – s = Select * from r Minus Select * from s;

A	B
X	1
Y	1

s - r

A	B
Y	3

- **Set Operators Properties:**
    - The union and the intersection are commutative operations

      $R \cup S = S \cup R$, and $R \cap S = S \cap R$
    - The union and the intersection are associative operations

      $R \cup (S \cup T) = (R \cup S) \cup T$, and $R \cap (S \cap T) = (R \cap S) \cap T$
    - The set difference operation is not commutative operation

      $R - S \neq S - R$

Cartesian (or Cross Product) Operator (×)

- If R and S are two relations, $R \times S$ is the set of all concatenated rows <x,y>, where x is a row in R and y is a row in S
    - *R* and *S* need not be type compatible
- $R \times S$ is expensive to compute:
    - Factor of two in the size of each row
    - Quadratic in the number of rows
- If R has n_R rows (denoted as $|R| = n_R$), and S has n_S rows, then R x S will have $n_R * n_S$ rows.
- $R \times S$ is equivalent to "Select * From R, S;"
    - Example:

r

A	B
X	1
Y	2

s

C	D	E
X	10	E1
Y	10	E1
Y	20	E2
Z	10	E2

$r \times s$ = Select * from r,s;

A	B	C	D	E
X	1	X	10	E1
X	1	Y	10	E1
X	1	Y	20	E2
X	1	Z	10	E2
Y	2	X	10	E1
Y	2	Y	10	E1
Y	2	Y	20	E2
Y	2	Z	10	E2

$\sigma_{A=C} (r \times s)$ = select * from r,s where A=C;

A	B	C	D	E

X	1	X	10	E1
Y	2	Y	10	E1
Y	2	Y	20	E2

- Generally, CROSS PRODUCT is not a meaningful operation
    - Can become meaningful when followed by other operations
    - Example (not meaningful):

Employee

SSN	FName	LName	Gender	SuperSSN	DNo

Dependent

ESSN	Dependent_Name	Gender	BDate	Relationship

$\text{Female\_Emps} \leftarrow \sigma_{\text{Gender='F'}}(\text{Employee})$
$\text{EmpNames} \leftarrow \pi_{\text{FNAME, LNAME, SSN}}(\text{Female\_Emps})$
$\text{Emp\_Dependents} \leftarrow \text{EmpNames x Dependent}$

-
    - Emp_Dependents will contain every combination of EmpNames and Dependent
    - whether or not they are actually related
    - To keep only combinations where the Dependent is related to the Employee, we add a SELECT operation as follows
      $\text{Actual\_Deps} \leftarrow \sigma_{\text{SSN=ESSN}}(\text{Emp\_Dependents})$
      $\text{Result} \leftarrow \pi_{\text{FNAME, LNAME, DEPENDENT\_NAME}}(\text{Actual\_Deps})$
    - Result will now contain the name of female employees and their dependents

- Example: Display employees names for employees who work in accounting department

Employee

SSN	Ename	Sal	Gender	SuperSSN	DNo

Department

DNo	DName	Location

$\pi_{\text{ENAME}}(\sigma_{\text{employee.dno=department.dno}}(\sigma_{\text{dname='Accounting'}}(\text{Employee} \times \text{Department})))$ =

Select Ename from Employee, Department
where employee.dno = department.dno and dname = ‘Accounting’;

- Example: find the names of all customers who live on the same street and in the same city as smith

Customer

SSN	Cname	City	Street

$\pi_{\text{CName}}(\sigma_{\text{street = s} \wedge \text{city = c}}($
$\text{Customer} \times (\rho_{\text{smith\_add(s,c)}}(\pi_{\text{street, city}}(\sigma_{\text{CName = 'Smith'}}(\text{Customer})))))$ =

Select c1.CName from customer c1, customer c2
where c2.CName='Smith' and c1.city = c2.city and c1.street = c2.street;

- Example: find the largest account balance in the bank

Account

AccNo	Balance	Date

π_{Balance} (Account) –
$\pi_{\text{Account.Balance}}$ ($\sigma_{\text{account.balance < d.balance}}$ (Account × (ρ_d (Account))))

Join Operator

- The JOIN operation is used to combine related rows from two relations into a single row.

١. **Theta Join Operator** ($\bowtie_\theta$)

- The Theta-Join is a specialized product containing only pairs that match on a supplied condition called join-condition.
- A theta join of R and S is the expression
  R $\bowtie_\theta$ S
  where θ is a conjunction of terms: A_i oper B_i
  in which A_i is an attribute of R; B_i is an attribute of S; and oper is one of =, <, >, ≥ ≠, ≤.
- R $\bowtie_\theta$ S = σ_θ (R × S) = Select * from R,S where θ;

- Example: **R**

A	B	C	D
a	1	x	4
b	2	y	5
c	4	z	4
d	8	x	5
e	1	y	4

S

E	F	G
5	a	x
4	b	y
3	c	y
2	a	x

R $\bowtie_{A<>\text{'a'} \wedge D<E}$ **S**

A	B	C	D	E	F	G
c	4	z	4	5	a	x
e	1	y	4	5	a	x

R ⋈$_{B = E}$ S (Equi_Join)

A	B	C	D	E	F	G
b	2	y	5	2	a	x
c	4	z	4	4	b	y

- Example: Display the names of all employees who earn more than their managers.

  Employee

ID	Ename	Salary	MgrId

  Manager

ID	MName	Salary

  π_{EName} (Employee ⋈ $_{MgrId=Manager.Id \text{ AND } Employee.Salary>Manager.Salary}$ Manager)

٢. **Natural Join Operator** (⋈)

- Special case of Equi_Join.
- Natural join requires that the two join attributes, or each pair of corresponding join attributes, have the same name in both relations. If this is not the case, a renaming operation is applied first.
- Natural join removes duplicate attributes.
- r ⋈ s = $\pi_{\text{Attribute\_List}}$ ($\sigma_{\text{Join\_Condition}}$ (r × s))

  where

  Attribute_List = *attributes* (r) ∪ *attributes* (s)

  (duplicates are eliminated) and *Join-Condition* has the form:

  $A_1 = A_1$ AND … AND $A_n = A_n$

  where $\{A_1 \ldots A_n\} = attributes(r) \cap attributes(s)$

- Note: let r(R) and s(S) be relations without any attributes in common; that is,

  R ∩ S = □. Then r ⋈ s = r × s.

- Example: **R**

A	B	C	D
a	1	x	4
b	2	y	5
c	4	z	4
d	8	x	5
e	1	y	4

S

E	F	G
5	a	x
4	b	y
3	c	y
2	a	x

R ⋈$\rho_{(B, F, G)}$ **(S) =** Select R.*, F, G from R, S where B=E;

A	B	C	D	F	G
b	2	y	5	a	x
c	4	z	4	b	y

- Example: **R**

A	B	C	D
x	1	x	a
y	2	z	a
z	4	y	b
x	1	z	a
w	2	y	b

S

B	D	E
1	a	x
3	a	y
1	a	z
2	b	w
3	b	e

$$\mathbf{R \bowtie S = \pi_{R.A, R.B, R.C, R.D, S.E} (\sigma_{R.B = S.B \wedge R.D = S.D} (R \times S))}$$

A	B	C	D	E
x	1	x	a	x
x	1	x	a	z
x	1	z	a	x
x	1	z	a	z
w	2	y	b	w

Complete Set of Relational Operations

- The set of operations including $\boldsymbol{\pi}$ (Projection), $\boldsymbol{\sigma}$ (Selection), **–** (Difference), $\boldsymbol{\rho}$ (Rename), (Union) and × (Cartesian Product) is called a complete set, because any other relational algebra expression can be expressed by combination of these five operations.
  - $R \cap S = (R \cup S) - ((R - S) \cup (S - R))$
  - $R \cap S = R - (R - S)$
  - $R \bowtie_{Condition} S = \sigma_{Condition} (R \times S)$

Examples of Queries in Relational Algebra

8.4 Examples of Queries in Relational Algebra

Banking Example:

Branch (branch_name, branch_city)

Customer (customer_name, customer_city, customer_street)

Account (account_number, branch_name, balance)

Loan (loan_number, branch_name, amount)

Depositor (customer_name, account_number)

Borrower (customer_name, loan_number)

Q1: Find all loans of over $1200.

Q2: Find the loan number for each loan of an amount greater than $1200.

Q3: Find the names of all customers who have a loan, an account, or both from the bank.

Q4: Find the names of all customers who have a loan and an account at bank.

Q5: Find the names of all customers who have a loan at the Perryridge branch.

Q6: Find the names of all customers who have a loan at the Perryridge branch but do not have an account at any branch of the bank.

Q7: Find the names of all branches with customers who have an account in the bank and who live in Harrison.

Q8: Find all customers who have an account from at least the “Downtown” and the “uptown” branches.

Company Example:

Employee (fname, minit, lname, SSN, address, sex, salary, superSSN, DNo)

Department (Dname, Dnumber, MGRSSN, MGRStartDate)

Dept_Locations (DNumber, DLocation)

Project (PName, PNumber, PLocation, DNum)

Works_On (ESSN, PNo, Hours)

Dependent (ESSN, Dependent_Name, Sex, BDate,Relationship)

Q1: Retrieve the name and address of all employees who work for the 'Research' department.

Q2: for every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birthdate.

Q3: make a list of project numbers for projects that involve an employee whose last name is 'smith', either as a worker or as a manager of the department that controls the project.

Q4: Retrieve the names of employees who have no dependents.

Q5: List the names of managers who have at least one dependent.